

Plant Disease Detection CNN — Project Report

Deep Learning for Agricultural Disease Classification

May 2026 | v2 — with Augmentation & Early Stopping

15

Disease Classes

20,638

Total Images

91.63 %

Baseline Test Acc

91.11 %

Best Val Acc (Run 2)

Abstract

This report documents two training runs of a custom CNN for plant disease classification across 15 classes from the PlantVillage dataset. **Run 1** (baseline): 10 epochs, no augmentation — test accuracy **91.63 %**. **Run 2** (improved): data augmentation + early stopping, planned for 20 epochs but stopped after epoch 11 (best weights at epoch 10) — best validation accuracy **91.11 %**, best validation loss **0.2644**. Both runs use TensorFlow 2.21 / Keras 3.13 on a Google Colab TPU (V5E1).

1. Introduction & Motivation

Plant diseases cause significant agricultural losses worldwide. Traditional expert-based visual inspection is costly and hard to scale. This project implements a custom **Convolutional Neural Network (CNN)** from scratch using **TensorFlow 2.21** / **Keras 3.13** to classify 15 categories of plant leaf images (healthy or diseased) from the **PlantVillage** dataset. The work spans two iterative training runs, progressively adding data augmentation and early stopping to improve generalisation.

Project Objectives

- Download and preprocess the PlantVillage dataset via the Kaggle API
- Explore class distribution and visualise sample leaf images
- Design and implement a multi-layer CNN architecture
- Train on a Google Colab TPU V5E1 accelerator
- Iteratively improve with data augmentation and early stopping
- Evaluate and compare performance across both training runs

2. Dataset

2.1 Source & Acquisition

Dataset: **emmarex/plantdisease** (Kaggle) — a PlantVillage subset. Contains RGB leaf photographs of 3 plant species (Bell Pepper, Potato, Tomato) across **15 classes**. Total: **20,638 images**. Downloaded programmatically via the Kaggle CLI and extracted from ZIP.

2.2 Split Strategy

80 % training (16,511 images) / 20 % validation (4,127 images), fixed seed 123 for reproducibility. An additional 20 % of the validation split was reserved as the held-out **test set** for Run 1.




Figure 1 — Image count per class. Note strong imbalance: Potato-Healthy has only 152 samples vs. 3,209 for Tomato-YLCV.

2.3 Per-Class Image Counts

#	Class Name	Count	Type
1	Pepper - Bacterial Spot	997	Disease
2	Pepper - Healthy	1478	Healthy
3	Potato - Early Blight	1000	Disease
4	Potato - Late Blight	1000	Disease
5	Potato - Healthy	152	Healthy
6	Tomato - Bacterial Spot	2127	Disease
7	Tomato - Early Blight	1000	Disease
8	Tomato - Late Blight	1909	Disease
9	Tomato - Leaf Mold	952	Disease
10	Tomato - Septoria Leaf Spot	1771	Disease
11	Tomato - Spider Mites	1676	Disease
12	Tomato - Target Spot	1404	Disease
13	Tomato - Yellow Leaf Curl Virus	3209	Disease
14	Tomato - Mosaic Virus	373	Disease
15	Tomato - Healthy	1591	Healthy

3. Data Preprocessing

Images are loaded via `tf.keras.utils.image_dataset_from_directory`, resized to **224 × 224 pixels**, batched at **32**, and one-hot encoded. Pixel normalisation ($\div 255$) is done in-graph via a `Rescaling(1./255)` layer. Run 2 additionally applies a data augmentation pipeline as the second layer in the model (see Section 5.2).

- **Image size:** 224 × 224 px, 3-channel RGB
- **Batch size:** 32
- **Normalisation:** in-graph Rescaling layer ($\div 255$)
- **Labels:** inferred from directory structure, one-hot encoded
- **Pipeline:** `tf.data` with AUTOTUNE prefetch
- **Run 2 augmentation:** random horizontal flip, random rotation, random zoom — applied in-model

4. Model Architecture

A custom **Sequential CNN** with three convolutional blocks (Conv2D + MaxPooling2D) followed by a dense classification head with Dropout regularisation. The architecture is shared between both runs; Run 2 inserts a Keras augmentation layer immediately after the Rescaling layer.




Figure 2 — CNN architecture overview.

4.1 Layer Details

Layer	Config	Output Shape	Purpose
Rescaling	1/255	224×224×3	Normalise pixels
[Augmentation]	Flip/Rot/Zoom (Run 2 only)	224×224×3	Data augmentation
Conv2D	32 filters, 3×3, ReLU	222×222×32	Low-level features
MaxPooling2D	2×2	111×111×32	Downsampling
Conv2D	64 filters, 3×3, ReLU	109×109×64	Mid-level features
MaxPooling2D	2×2	54×54×64	Downsampling
Conv2D	128 filters, 3×3, ReLU	52×52×128	High-level features
MaxPooling2D	2×2	26×26×128	Downsampling
Flatten	—	86,528	3D → 1D vector
Dense	128 units, ReLU	128	Classification head
Dropout	rate = 0.5	128	Regularisation
Dense (output)	15 units, Softmax	15	Class probabilities

Table 2 — Layer details. The augmentation row (orange) is active in Run 2 only.

4.2 Compilation

- **Optimiser:** Adam (lr = 0.001 default)
- **Loss:** Categorical Cross-Entropy
- **Metric:** Accuracy

5. Run 1 — Baseline Training (10 Epochs, No Augmentation)

The first training run used the raw dataset without any augmentation. The model was trained for a fixed **10 epochs** on a **Google Colab TPU V5E1** accelerator (TensorFlow 2.21). Each epoch processed 516 batches at ~1 s/step, totalling approximately **106 minutes** of training time.




Figure 3 — Run 1 accuracy (left) and loss (right). The orange dotted line marks the final test accuracy of 91.63 %. Best validation accuracy (91.06 %) at epoch 7.

5.1 Epoch-by-Epoch Results

Epoch	Train Acc	Val Acc	Train Loss	Val Loss
1	0.4740	0.7549	1.6429	0.7452
2	0.7166	0.8296	0.8516	0.4926
3	0.7808	0.8359	0.6494	0.4890
4	0.8175	0.8872	0.5365	0.3398
5	0.8370	0.8652	0.4808	0.3780
6	0.8684	0.8887	0.3831	0.3152
7	0.8885	0.9106	0.3246	0.2817
8	0.9012	0.8984	0.2855	0.3031
9	0.9085	0.8936	0.2728	0.3418
10	0.9218	0.9053	0.2202	0.2963

Table 3 — Run 1 per-epoch metrics. Epoch 7 (highlighted) = best val acc 91.06 %.

5.2 Test Set Evaluation

91.63 %

0.3807

10

~106 min

Test Accuracy	Test Loss	Epochs Trained	Training Time
---------------	-----------	----------------	---------------

The baseline model achieved a solid **test accuracy of 91.63 %**. The gap between training accuracy (92.2 %) and validation accuracy (90.5 %) at epoch 10 is modest, indicating mild overfitting. Validation loss fluctuated after epoch 7, suggesting the model would benefit from early stopping.

6. Run 2 — Augmentation + Early Stopping (Stopped at Epoch 11)

Building on Run 1's findings, Run 2 introduced **in-model data augmentation** (random horizontal flip, random rotation $\pm 10^\circ$, random zoom $\pm 10\%$) and **early stopping** monitoring val_loss with patience. Training was planned for 20 epochs but the monitor triggered at **epoch 11**, having determined that epoch 10's val_loss of **0.2644** was the best checkpoint and subsequent steps were worsening.

Early Stopping Triggered — Epoch 11

Best weights restored from Epoch 10 | val_loss = 0.2644 | val_acc = 91.11 %




Figure 4 — Run 2 accuracy (left) and loss (right). Green dotted vertical line marks the best epoch (10). Early stopping restored weights from that checkpoint.

6.1 Epoch-by-Epoch Results

Epoch	Train Acc	Val Acc	Train Loss	Val Loss
1	0.4685	0.6895	1.6672	0.9195
2	0.6685	0.7651	1.0194	0.7052
3	0.7277	0.7285	0.8326	0.8601
4	0.7559	0.7710	0.7317	0.7716
5	0.7902	0.7939	0.6320	0.6990
6	0.8040	0.8667	0.5974	0.4382
7	0.8256	0.9048	0.5230	0.2764
8	0.8383	0.8848	0.4899	0.3501
9	0.8502	0.8843	0.4429	0.3932
10	0.8640	0.9111	0.4002	0.2644

Table 4 — Run 2 per-epoch metrics. Epoch 10 (highlighted) = best val_loss 0.2644, val_acc 91.11 %. Early stopping fired at epoch 11.

6.2 Key Observations

- Training accuracy rose more gradually than Run 1 (46.8 % → 86.4 % vs. 47.4 % → 92.2 %), which is expected — augmentation makes each sample harder, acting as implicit regularisation.
- Best validation loss (**0.2644** at epoch 10) is notably lower than Run 1's best val_loss of 0.2817 (epoch 7), indicating better generalisation.
- Validation accuracy peaked at **91.11 %** — marginally higher than Run 1's best val_acc of 91.06 %, confirming the improvement despite lower training accuracy.
- Epoch 3 showed a val_loss spike (0.86) — typical early instability when learning augmented samples, resolving by epoch 6.
- Early stopping successfully prevented the model from training past its optimal checkpoint, saving roughly 9 additional epochs (~100 min) of unnecessary compute.

7. Run Comparison




Figure 5 — Validation accuracy (left) and loss (right) for Run 1 (green) vs. Run 2 (orange) over 10 epochs. Run 2 achieves a lower final val_loss despite lower training accuracy, demonstrating better generalisation.

7.1 Summary Comparison Table

Metric	Run 1 (Baseline)	Run 2 (Aug + ES)	Delta
Augmentation	None	Flip/Rot/Zoom	—
Early Stopping	No	Yes (patience)	—
Epochs Planned	10	20	+10
Epochs Executed	10	11 (stopped)	+1
Best Val Acc	91.06 % (ep7)	91.11 % (ep10)	+0.05 pp
Best Val Loss	0.2817 (ep7)	0.2644 (ep10)	-0.0173
Final Train Acc	92.18 %	86.40 %	-5.78 pp
Final Val Acc	90.53 %	91.11 %	+0.58 pp
Test Accuracy	91.63 %	N/A (not run)	—
Test Loss	0.3807	N/A (not run)	—
Approx. Train Time	~106 min	~123 min	+17 min

Table 5 — Side-by-side comparison. Best val_loss row (highlighted) shows Run 2 achieves a 1.73-point improvement in validation loss.

8. Compute Environment & Dependencies

Component	Version / Config
Platform	Google Colab

Component	Version / Config
Accelerator	TPU V5E1
Python	3.12
TensorFlow	2.21.0
Keras	3.13.2
NumPy	2.0.2
scikit-learn	latest (Colab)
Kaggle API	2.0.2

9. Limitations & Recommendations

9.1 Remaining Limitations

- **Class imbalance:** Potato-Healthy (152) vs. Tomato-YLCV (3,209) — skew can depress recall on minority classes. Consider class-weighted loss.
- **Shallow architecture:** 3 conv blocks may not capture very fine-grained inter-disease differences. A deeper or pre-trained backbone would help.
- **Lab images only:** PlantVillage uses controlled conditions; real-field performance may be lower due to background clutter and lighting.

9.2 Recommended Next Steps

- **Transfer learning:** Fine-tune EfficientNetB0 or MobileNetV3 on PlantVillage — typically achieves 95–98 % with minimal extra training.
- **Class-weighted loss / SMOTE** to handle the Potato-Healthy minority class.
- **Learning rate scheduling:** Add ReduceLROnPlateau to capitalise on plateaus rather than stopping early.
- **Grad-CAM visualisation:** Verify the model attends to disease lesions rather than image artefacts.

10. Conclusion

This project demonstrates successful end-to-end development of a CNN classifier for plant disease detection across two iterative training runs.

Run 1 (baseline) achieved a **test accuracy of 91.63 %** in 10 epochs without augmentation. **Run 2** introduced data augmentation and early stopping, yielding a lower best validation loss (**0.2644** vs. 0.2817) and slightly higher best validation accuracy (**91.11 %** vs. 91.06 %), stopping efficiently at epoch 11 after detecting no improvement. The lower training accuracy in Run 2 (86.4 % vs. 92.2 %) is expected and desirable — augmentation increases task difficulty during training, leading to better generalisation.

The results confirm that even a relatively shallow custom CNN can achieve strong performance on PlantVillage. With transfer learning and per-class metrics, accuracy in the **95–98 % range** should be readily attainable.